

DogFit: Domain-guided Fine-tuning for Efficient Transfer Learning of Diffusion Models



Yara Bahram, Mohammad Shateri, Eric Granger
yara.mohammadi-bahram@livia.etsmtl.ca

The 40th Annual AAAI Conference on Artificial Intelligence
JANUARY 20 – JANUARY 27, 2026 | SINGAPORE



Main Contributions

- We propose domain-guided fine-tuning (**DogFit**), a controllable guidance mechanism for diffusion transfer learning that improves target domain alignment without needing double forward passes at test-time.
- We show that during fine-tuning, the source model offers stronger marginal estimates than the learning target model, making it better suited for generating guidance signals.
- For transfer learning across six target datasets on DiT and SiT models, **DogFit** outperforms state-of-the-art guidance methods in majority of cases while being up to 2x faster.

Motivation and Method


Diffusion Transfer Learning is Still Challenging

- **Fine-tuning** on small target datasets generalizes poorly and suffers from **low fidelity**.
- **Test-time guidances** like classifier-free guidance (CFG)^[1] and domain guidance (DoG)^[2] improve fidelity, but sampling needs **double forward passes**.
- **Training-time guidances** like model guidance (MG)^[3] do not require double forward passes, but exhibit **bad target alignment** and **no test-time guidance control**.


DogFit to the Rescue

- Improved fidelity
- Target domain alignment
- Single forward pass sampling
- Test-time guidance control

DogFit Tweaks the Denoising Training Loss



Adding Test-time Guidance Control



- Sample from shifted exponentially decaying distribution (SEDD):

$$w = 1 + z, \quad z \sim \mathcal{P}(z),$$

$$\mathcal{P}(z) = \lambda e^{-\lambda z}, \quad z \geq 0,$$

A plot of the probability density function $\mathcal{P}(z)$ versus z shows a curve starting at $z=0$ and decaying towards zero as z increases.

DogFit and Friends for Guiding Generation in Transfer Learning




Figure 1: Transfer learning via guidance mechanisms for synthetic data created using a mixture of Gaussian distributions. (a) The red region defines the target domain, while the blue background represents the source distribution. (b) CFG and MG prioritize class separability without considering the source distribution, often pushing samples toward out-of-distribution areas in the target domain. (c) **DogFit** and DoG utilize source information to emphasize movement towards the core of the target manifold, improving domain alignment. (*) Sampling-time guidance methods (DoG and CFG) operate by computing the guidance offset, whereas training-time guidance methods (**DogFit** and MG) learn the guided direction directly.

Fast Sampling, Simple Training, and Guidance Control




Figure 2: Illustration of DogFit during training and sampling : (a-b) DogFit internalizes guidance behavior, removing the need for inference-time double forward passes. (c) During training, the guidance value is treated as an input, allowing inference-time control. The simple case of our method, with fixed guidance, simply requires removing the w and 1 from the target model conditions. (d) w is embedded it as an extra input during fine-tuning, allowing inference-time control. (*) Dashed lines denote paths that do not propagate gradients.

Results

Controllable Guidance Knob at Test-time




Figure 4: Effect of controllable guidance at test time for the Food and ArtBench domains. (Left) FID behavior across test-time w s with different training-time λ s. (Right) Corresponding generated samples for fixed $\lambda = 3$ (95% of sampled w values in $[1, 2]$), across varying test-time w . Gray cells indicate extreme FID values (≈ 350) resulting from collapsed generations.

When and Where to Guide Matters




Figure 5: Ablation study on guidance schedules in DogFit. (a) Varying the late-start threshold τ_s to control when guidance begins. (b) Varying the cut-off threshold τ_c to restrict guidance to later denoising steps. Performed using FID on DiT-XL2.

# Steps	Food FID↓	Art FID↓	Train FID↓
MG	60.60	57.91	108.26
DogFit	21.13	18.97	39.80
MG + DogFit	11.13	10.64	19.91
MG	7.60	8.09	13.38
DogFit	6.75	7.00	12.70

Table 1: Varying the number of sampling steps for MG and DogFit.

Variant	Food FID↓	Art FID↓	Train Params (M)
DiT-XL2	10.64	16.32	675.42 (100%)
DogFit + DiT-XL2	11.86	15.80	0.75 (0.11%)

Table 2: Using DogFit with Parameter-efficient Fine-tuning (PEFT).

References

- Ho, Jonathan, and Tim Salimans. "Classifier-free diffusion guidance." *arXiv preprint* (2022).
- Zhong, Jincheng, et al. "Domain guidance: A simple transfer approach for a pre-trained diffusion model." *ICLR* 2025.
- Tang, Zhicong, et al. "Diffusion models without classifier-free guidance." *arXiv preprint* (2025).
- Xie, Enze, et al. "Diffit: Unlocking transferability of large diffusion models via simple parameter-efficient fine-tuning." *ICCV* 2023.